

Linear Regression

The purpose of these coding vignettes is to explore how Linear Regression for classification works. As with the Perceptron Learning Algorithm implementation, I have created my own target function f and data set \mathcal{D} . We take $d = 2$ so we can visualize the problem, and assume $\mathcal{X} = [-1, 1] \times [-1, 1]$ with uniform probability of picking each $\mathbf{x} \in \mathcal{X}$. In each run, we choose a random line in the plane as our target function f (we do this by taking two random, uniformly distributed points in $[-1, 1] \times [-1, 1]$ and taking the line passing through them), where one side of the line maps to $+1$ and the other maps to -1 . Then we choose the inputs \mathbf{x}_n of the data set as random points (uniformly in \mathcal{X}), and evaluate the target function on each \mathbf{x}_n to get the corresponding output y_n .

```
(* Clear globals *)
Clear[GenerateX, DoLinearRegressionExperiment, Ein,
  GetYInterceptForm, LinearTarget, NoFeature, RiskNoiseFlip]

(* Generate our  $\mathcal{X} = [-1, 1] \times [-1, 1]$  space with  $d$  uniformly distributed
  points and a random  $\mathcal{X}$ -partitioning solution function  $f$ . *)

GenerateX[OptionsPattern[]] :=
Module[{
  D, f, d = 2,
  f1, f2,
  m, a, b, c,
  t1, t2,
  dotTest1, dotTest2, DotTest
},
  f1 = RandomReal[{-1, 1}, {1, d}][[1]];
  f2 = RandomReal[{-1, 1}, {1, d}][[1]];

  m =  $\frac{f1[[2]] - f2[[2]]}{f1[[1]] - f2[[1]]}$ ;
  a = -m;
  b = 1;
  c = m f1[[1]] - f1[[2]];
  f = {{c}, {a}, {b}};

  (* f should not dot to zero for our two original points! *)
  t1 = {1, f1[[1]], f1[[2]]};
  t2 = {1, f2[[1]], f2[[2]]};

  DotTest[v_, s_] := If[Abs[v] > 0.00000000001, Throw[s], 0];
  DotTest[t1.f, "t1 dot test failed"];
  DotTest[t2.f, "t2 dot test failed"];

  D = RandomReal[{-1, 1}, {OptionValue[DSize], d}];

  {D, f}
]

Options[GenerateX] = {DSize -> 100};
```

```

LinearTarget[f_, X_] := Sign[X.f];

NoFeature[X_] := X;

RiskNoiseFlip[percent_] := If[RandomReal[] ≤  $\frac{\text{percent}}{100}$ , -1, 1];

DoLinearRegressionExperiment[X_, OptionsPattern[]] :=
Module[{
  D, f, y,
  X, Xf, Xfdag,
  w
},
  {D, f} = X;
  X = Function[x, Prepend[x, 1]] /@ D;
  Xf = OptionValue[DataFeature][X];
  y = RiskNoiseFlip[OptionValue[Noise]] * OptionValue[TargetFunction][f, X];
  Xfdag = PseudoInverse[Xf];
  w = Xfdag.y;
  {w, X, y, D, f}
]
Options[DoLinearRegressionExperiment] =
  {TargetFunction → LinearTarget, DataFeature → NoFeature, Noise → 0};

ClassificationEin[X_, w_, y_, OptionsPattern[]] :=
Module[{
  N, misses, sumOfMisses
},
  N = Length[y];
  (*misses=MapThread[(If[#1≠#2,1,0])&,{Sign[X.w],y}];*)
  misses =
    MapThread[(If[#1 ≠ #2, 1, 0]) &, {OptionValue[TargetFunction][w, X], y}];
  sumOfMisses = Total[misses];
   $\frac{1}{N}$  sumOfMisses
]
Options[ClassificationEin] = {TargetFunction → LinearTarget};

RegressionsEin[X_, w_, y_] :=
Module[{
  N
},
  N = Length[y];
   $\frac{1}{N}$  Norm[Sign[X.w] - y]2
]

(*
AltEin[X_, w_, y_] :=
Module[{

```

```

      N
    },
    N=Length[y];
    
$$\frac{1}{N} \sum_{n=1}^N (\text{Transpose}[w] \cdot X[[n]] - y[[n]])^2$$

  ]
*)

GetYInterceptForm[w_] :=
Module[{
  a, b, c, M, B
},
  a = w[[2]];
  b = w[[3]];
  c = w[[1]];
  M = - $\frac{a}{b}$ ; (*If[b!=0, - $\frac{a}{b}$ , 10000];*)
  B = - $\frac{c}{b}$ ; (*If[b!=0, - $\frac{c}{b}$ , 10000];*)
  {M, B}
]

(* I don't think we're going to need this,
but who knows so I'll keep it around until we
get the perceptron reimplemented in Mathematica. *)
(*
GetBoundaryPoints[w_] :=
Module[{
  a, b, c,
  xAtYmax, xAtYmin,
  yAtXmax, yAtXmin,
  m
},
  a = w[[2]];
  b = w[[3]];
  c = w[[1]];
  xAtYmax =  $\frac{-c-b}{a}$ ;
  xAtYmin =  $\frac{-c+b}{a}$ ;
  yAtXmax =  $\frac{-c-a}{b}$ ;
  yAtXmin =  $\frac{-c+a}{b}$ ;
  m = Sign[- $\frac{a}{b}$ ];
  {x1, y1, x2, y2} =
  Switch[
    m > 0, True, Switch[
      ] *)

```

Finding g when N = 100

Take $N = 100$. Use Linear Regression to find g and evaluate E_{in} , the fraction of in-sample points which got classified incorrectly. Repeat the experiment 1000 times and take the average (keep the g 's as they will be used again later). Note the average E_{in} .

```
(* Clear globals *)
Clear[Experiment1]
Clear[elavEin, elfs, elgs, elD, elruns, n]
Clear[DoInteractiveLR, DoAutoLR]

(* The interactive mode gives us the option of stepping through
   randomly generated data/f/g just to see what they look like. *)
DoInteractiveLR[] :=
DynamicModule[{
  runs = 1000, N = 100,
  g, gs,
  X, Y,
  D, f,
  avEin,
  i,
  fm, fb,
  sm, sb
},
gs = {};
i = 1;
{g, X, Y, D, f} = DoLinearRegressionExperiment[GenerateX[]];
EventHandler[Dynamic[
  (*avEin=Mean[Function[x,Ein[X,x,Y]]/@gs];*)
  avEin = ClassificationEin[X, g, Y];
  {fm, fb} = GetYInterceptForm[f];
  {sm, sb} = GetYInterceptForm[g];
  Show[ListPlot[D, PlotLabel -> StringForm[
    "Linear regression results\n(runs=`, N=`) : \naverage (class.) Ein=``",
    runs, N, avEin]],
    Plot[{fm x + fb + 0.05, sm x + sb}, {x, -2, 2}]
  ]
],
{"KeyDown" -> ({g, X, Y, D, f} = DoLinearRegressionExperiment[GenerateX[]]; i++)}
]

(* Auto mode implements what the problem asks for:
   it runs the tests and then outputs a fomatted string bearing the results. *)
DoAutoLR[] :=
Module[{
  runs = 1000, N = 100,
  g, gs,
  Y, Ys,
  X, Xs,
  f, fs,
  D,
  avEin,
  i, n
},
],
fs = {};
```

```

gs = {};
ys = {};
Xs = {};
For[i = 1, i ≤ runs, i++, (
  {g, X, y, D, f} = DoLinearRegressionExperiment[GenerateX[DSize → 100]];
  AppendTo[fs, f];
  AppendTo[gs, g];
  AppendTo[ys, y];
  AppendTo[Xs, X]
)];
avEin =
  Mean[MapThread[Function[{x, y, z}, ClassificationEin[x, y, z]], {Xs, gs, ys}]];
n = Length[D];
{fs, gs, D, avEin, runs, n}
]

Experiment1[interactive_] := If[interactive, DoInteractiveLR[], DoAutoLR[]]
{elfs, elgs, elD, elavEin, elruns, eln} = Experiment1[False];
(*Experiment1[True]*)
StringForm[
  "Linear regression results\n(runs=`, N=`) : \naverage (class.) Ein=`",
  elruns, eln, elavEin × 1.]

Linear regression results
(runs=1000, N=100):
average (class.) Ein=0.03898`

```

Evaluating E_{out} with $N = 1000$

Now, generate 1000 fresh points and use them to estimate the out-of-sample error E_{out} of g that we got above (number of misclassified out-of-sample points / total number of out-of-sample points). Again, we run the experiment 1000 times and take the average and note the average E_{out} .

```

(* clear globals *)
Clear[Experiment2]
Clear[e2fs, e2gs, e2avEout, e2runs, e2n]

```

```

DoEoutTest[fs_, gs_] :=
Module[{
  runs = 1000, N = 1000,
  g,
  y, ys,
  X, Xs,
  f,
  D,
  avEout,
  i
},
ys = {};
Xs = {};
For[i = 1, i ≤ runs, i++, (
  {g, X, y, D, f} =
    DoLinearRegressionExperiment[{GenerateX[DSize → N][[1]], fs[[i]]}];
  AppendTo[ys, y];
  AppendTo[Xs, X];
)];
avEout =
  Mean[MapThread[Function[{x, y, z}, ClassificationEin[x, y, z]], {Xs, gs, ys}]];
n = Length[D];
{fs, gs, avEout, runs, n}
]

Experiment2[] := DoEoutTest[e1fs, e1gs];
{e2fs, e2gs, e2avEout, e2runs, e2n} = Experiment2[];
StringForm[
  "Linear regression results\n(runs=`, N=`):\naverage (class.) Eout=`",
  e2runs, e2n, e2avEout × 1.]

Linear regression results
(runs=1000, N=1000):
average (class.) Eout=0.047891`

```

Using linear regression to determine initial weights for PLA, N=10

Using $N = 10$, we find the weights using Linear Regression and then use them as a vector of initial weights for the Perceptron Learning Algorithm. We then run PLA until it converges to a final vector of weights that completely separates all the in-sample points. We can then examine the average number of iterations (over 1000 runs) that PLA takes to converge.

```

(* clear globals *)
Clear[DoPLAExperiment, DoInteractivePLAExperiment, DoLRtoPLASeries]
Clear[e3runs, e3n, e3AveIterations]
Clear[Experiment3]

```

```

DoPLAStep[w_, f_, D_, Converged_, Misses_] :=
Module[{
  x, h, chk,
  missedPoints,
  misses,

```

```

numSamples,
converged,
w,
i, ri
},
converged = Converged;
misses = Misses;
w = w;
If[converged == 0, (
  converged = 1;
  missedPoints = {};
  numSamples = Length[D];

  For[i = 1, i ≤ numSamples, i++, (
    (* cycle through the samples looking for misses *)
    x = Transpose[ArrayReshape[Prepend[D[[i]], 1], {1, 3}]];
    h = Sign[(Transpose[w].x) [[1, 1]]];
    chk = Sign[(Transpose[f].x) [[1, 1]]];

    If[chk ≠ h, (
      (* if we have a miss...*)
      converged = 0;
      AppendTo[missedPoints, x];
      (*iterMisses++;*)
      0), 0]
    )];

  If[converged == 0, (
    (* if we've missed any points,
    then pick one of them at random for the hypothesis update *)
    misses++;
    ri = RandomInteger[{1, Length[missedPoints]}];
    x = missedPoints[[ri]];
    h = Sign[(Transpose[w].x) [[1, 1]]];
    chk = Sign[(Transpose[f].x) [[1, 1]]];
    w = w + chk * x; (* avoiding "cannot assign to raw object 0" *)
    (*w=q*) (* ...maybe because it was self-
    assigning and the last value in a "block" thing? *)
    (*ah no I think it's because it's declared as a parameter!*)
    )];
  )];
{w, converged, misses}
]

DoPLAExperiment[x_, OptionsPattern[]] :=
Module[{
  D, f, w, g,
  misses = 0,
  converged = 0
},
{D, f} = x;
w = OptionValue[w];

While[converged == 0, (
  {w, converged, misses} = DoPLAStep[w, f, D, converged, misses];

```

```

    ]];

    g = w;
    {g, Length[D], misses}
  ]
Options[DoPLAExperiment] = {w → {{0}, {0}, {0}}};

DoInteractivePLAExperiment[χ_] :=
  DynamicModule[{
    D, f, w, g,
    Dpos, Dneg,
    fm, fb, sm, sb,
    misses = 0,
    converged = 0,
    i = 0
  },
  {D, f} = χ;
  w = {{0.000000001}, {0.000000001}, {0.000000001}};
  {w, converged, misses} = DoPLAStep[w, f, D, converged, misses];

  EventHandler[Dynamic[{
    {fm, fb} = GetYInterceptForm[f];
    {sm, sb} = GetYInterceptForm[w];
    Dpos = Select[D, Sign[ArrayReshape[Prepend[#, 1], {1, 3}].f][[1, 1]] ≥ 0 &];
    Dneg = Select[D, Sign[ArrayReshape[Prepend[#, 1], {1, 3}].f][[1, 1]] < 0 &];
    Show[
      ListPlot[{Dpos, Dneg}, PlotMarkers → {"+", "-"}, PlotLabel → StringForm[
        "i=``, misses=``, i, misses]], Plot[{fm x + fb, sm x + sb}, {x, -2, 2}]
    ]
  }],
  {"KeyDown" ⇒ ({w, converged, misses} = If[converged == 0,
    DoPLAStep[w, f, D, converged, misses], {w, converged, misses}]; i++)}
  ]
]

DoLRtoPLASeries[runs_] :=
  Module[{
    g, X, y, D, f, i, n,
    iterations, iterationses, avIterations
  },
  iterationses = {};
  For[i = 1, i ≤ runs, i++, (
    {g, X, y, D, f} = DoLinearRegressionExperiment[GenerateX[DSize → 10]];
    {g, n, iterations} = DoPLAExperiment[{D, f}, w → g];
    AppendTo[iterationses, iterations];
  )];
  avIterations = Mean[iterationses];
  {runs, Length[D], avIterations}
]

(* interactive version *)
(*Experiment3[] := DoInteractivePLAExperiment[GenerateX[]]
Experiment3[] *)

(* single-run test version *)

```



```
(*Experiment3[]:=DoPLAExperiment[GenerateX[]];
{e3g,e3n,e3iterations}=Experiment3[];
StringForm["PLA results\n(N=``):\niterations to convergence=``",
  e3n, e3iterations]*)

(* linear regression to PLA version *)
Experiment3[]:=DoLRtoPLASeries[1000];
{e3runs,e3n,e3AveIterations}=Experiment3[];
StringForm[
  "LR to PLA results\n(runs=``, N=``):\naverage iterations to convergence=``",
  e3runs, e3n, e3AveIterations*1.]
```

```
LR to PLA results
(runs=1000, N=10):
average iterations to convergence=4.005`
```

Nonlinear Transformation

In these experiments we again apply Linear Regression for classification. Using the target function:

$$f(x_1, x_2) = \text{sign}(x_1^2 + x_2^2 - 0.6)$$

...we generate a training set of $N = 1000$ points on $\mathcal{X} = [-1, 1] \times [-1, 1]$ with a uniform probability of picking each $\mathbf{x} \in \mathcal{X}$ and generate simulated noise by flipping the sign of the output in a randomly selected 10% subset of the generated training set.

Linear Regression (sans transformation)

We carry out Linear Regression without transformation, i.e., with feature vector:

$$(1, x_1, x_2),$$

to find the weight \mathbf{w} . We will determine the classification in-sample error E_{in} by running the experiment 1000 times and taking the average E_{in} to reduce variation in our results.

```
(* clear globals *)
Clear[NonlinearTarget]
Clear[Experiment4, DoNonlinearDataSeries]
Clear[{e4fs, e4gs, e4D, e4avEin, e4avG, e4runs, e4n}]
```

```

NonlinearTarget[f_, X_] := Sign[(#2^2 + #3^2 - 0.6)] & @@@ X;

DoNonlinearDataSeries[runs_, N_] :=
Module[{
  g, gs,
  y, ys,
  X, Xs,
  f, fs,
  D,
  avEin, avG,
  i, n
},
  fs = {};
  gs = {};
  ys = {};
  Xs = {};
  For[i = 1, i ≤ runs, i++, (
    {g, X, y, D, f} = DoLinearRegressionExperiment[
      GenerateX[DSize → N], TargetFunction → NonlinearTarget, Noise → 10];
    AppendTo[fs, f];
    AppendTo[gs, g];
    AppendTo[ys, y];
    AppendTo[Xs, X]
  )];
  avEin =
    Mean[MapThread[Function[{x, y, z}, ClassificationEin[x, y, z]], {Xs, gs, ys}]];
  avG = Mean[gs];
  n = Length[D];
  {fs, gs, D, avEin, avG, runs, n}
]

Experiment4[] := DoNonlinearDataSeries[1000, 1000];
{e4fs, e4gs, e4D, e4avEin, e4avG, e4runs, e4n} = Experiment4[];
StringForm["Linear regression results\n(runs=`, N=`, Target=Nonlinear,
  Features=None):\ng=`` + ``x1 + ``x2\naverage (class.) Ein=``",
  e4runs, e4n, e4avG[[1]], e4avG[[2]], e4avG[[3]], e4avEin × 1.]

Linear regression results
(runs=1000, N=1000, Target=Nonlinear, Features=None):
g=0.04742016410294729` + 0.00325062 x1 + -0.00129323 x2
average (class.) Ein=0.507388`

```

Using nonlinear transformation, N = 1000

Now, we transform the N = 1000 training data into the following nonlinear feature vector:

$$(1, x_1, x_2, x_1 x_2, x_1^2, x_2^2)$$

and then find the vector $\tilde{\mathbf{w}}$ that corresponds to the solution of Linear Regression and display our final hypothesis, averaged over a few runs to make sure our answer is stable.

```
(* clear globals *)
Clear[NonlinearFeature]
Clear[Experiment5, DoNonlinearDataSeriesWithFeatures]
Clear[{e5fs, e5gs, e5D, e5avEin, e5avG, e5runs, e5n}]
```

```
NonlinearFeature[X_] := {1, #2, #3, #2 #3, #2^2, #3^2} & @@@ X;
```

```
DoNonlinearDataSeriesWithFeatures[runs_, N_, D_, fs_] :=
```

```
Module[{
  g, gs,
  y, ys,
  X, Xs,
  D, f,
  avEin, avG,
  i, n
},
  gs = {};
  ys = {};
  Xs = {};
  D = D;
  For[i = 1, i ≤ runs, i++, (
    {g, X, y, D, f} = DoLinearRegressionExperiment[{D, fs[[i]]}, TargetFunction →
      NonlinearTarget, Noise → 10, DataFeature → NonlinearFeature];
    AppendTo[gs, g];
    AppendTo[ys, y];
    AppendTo[Xs, X]
  )];
  avEin = Mean[MapThread[Function[{x, y, z}, ClassificationEin[
    x, y, z, TargetFunction → NonlinearTarget]], {Xs, gs, ys}]];
  avG = Mean[gs];
  n = Length[D];
  {fs, gs, D, avEin, avG, runs, n}
]
```

```
Experiment5[] := DoNonlinearDataSeriesWithFeatures[1000, 1000, e4D, e4fs];
```

```
{e5fs, e5gs, e5D, e5avEin, e5avG, e5runs, e5n} = Experiment5[];
```

```
StringForm["Linear regression results\n(runs=``,
  N=`, Target=Nonlinear, Features=Nonlinear):\ng=`` + ``x1
  + ``x2 + ``x1x2 + ``x12 + ``x22`\naverage (class.) Ein=``",
  e5runs, e5n, e5avG[[1]], e5avG[[2]], e5avG[[3]], e5avG[[4]],
  e5avG[[5]], e5avG[[6]], e5avEin × 1.]
```

Linear regression results

(runs=1000, N=1000, Target=Nonlinear, Features=Nonlinear):

$g = -0.95734 + -0.00543674 x_1 +$

$0.0180005 x_2 + 0.0173879 x_1 x_2 + 1.53297 x_1^2 + 1.50226 x_2^2$

average (class.) $E_{in} = 0.119$

E_{out} for nonlinear hypothesis

Estimate the value to the classification out-of-sample error E_{out} of our hypothesis above by generat-

ing a new set of 1000 points and adding noise, as before, then averaging over 1000 runs to reduce the variation in results.

```
(* clear globals *)
Clear[Experiment6, DoNonlinearDataSeriesEoutWithFeatures]
Clear[{e6fs, e6gs, e6D, e6avEin, e6avG, e6runs, e6n}]

DoNonlinearDataSeriesEoutWithFeatures[runs_, N_, D_, fs_, gs_] :=
Module[{
  g,
  y, ys,
  X, Xs,
  D, f,
  avEin, avG,
  i, n
},
  ys = {};
  Xs = {};
  D = D;

  For[i = 1, i ≤ runs, i++, (
    {g, X, y, D, f} = DoLinearRegressionExperiment[
      {GenerateX[DSize → N][[1]], fs[[i]]}, TargetFunction → NonlinearTarget,
      Noise → 10, DataFeature → NonlinearFeature];
    AppendTo[ys, y];
    AppendTo[Xs, X]
  )];
  avEout = Mean[MapThread[Function[{x, y, z}, ClassificationEin[
    x, y, z, TargetFunction → NonlinearTarget]], {Xs, gs, ys}]];
  avG = Mean[gs];
  n = Length[D];
  {fs, gs, D, avEout, avG, runs, n}
]

Experiment6[] :=
  DoNonlinearDataSeriesEoutWithFeatures[e5runs, e5n, e5D, e5fs, e5gs];
{e6fs, e6gs, e6D, e6avEin, e6avG, e6runs, e6n} = Experiment6[];
StringForm["Linear regression results\n(runs=``,
  N=`, Target=Nonlinear, Features=Nonlinear):\ng=`` + ``x1
  + ``x2 + ``x1x2 + ``x12 + ``x22`\naverage (class.) Eout=``",
  e6runs, e6n, e6avG[[1]], e6avG[[2]], e6avG[[3]], e6avG[[4]],
  e6avG[[5]], e6avG[[6]], e6avEin × 1.]
```

Linear regression results

```
(runs=1000, N=1000, Target=Nonlinear, Features=Nonlinear):
g=-0.95734 + -0.00543674 x1 +
  0.0180005 x2 + 0.0173879 x1x2 + 1.53297 x12 + 1.50226 x22
average (class.) Eout=0.085`
```